

We train an simple neural network to do task of image recognition, the data we used is CIFAR10 data set. Following the code.

```
In [1]: import torch
import torchvision
import torchvision.transforms as transforms
```

```
In [2]: #read data
transform = transforms.Compose(
    [transforms.ToTensor(),
     transforms.Normalize((0.5, 0.5, 0.5), (0.5, 0.5, 0.5))])

trainset = torchvision.datasets.CIFAR10(root='./data', train=True,
                                         download=True, transform=transform)
trainloader = torch.utils.data.DataLoader(trainset, batch_size=4,
                                           shuffle=True, num_workers=2)

testset = torchvision.datasets.CIFAR10(root='./data', train=False,
                                         download=True, transform=transform)
testloader = torch.utils.data.DataLoader(testset, batch_size=4,
                                          shuffle=False, num_workers=2)

classes = ('plane', 'car', 'bird', 'cat',
           'deer', 'dog', 'frog', 'horse', 'ship', 'truck')
```

Downloading <https://www.cs.toronto.edu/~kriz/cifar-10-python.tar.gz> to ./data/cifar-10-python.tar.gz

100%|████████████████████| 170498071/170498071 [00:06<00:00, 28396235.45 it/s]

Extracting ./data/cifar-10-python.tar.gz to ./data  
Files already downloaded and verified

```
In [3]: import torch.nn as nn
import torch.nn.functional as F
```

```
In [4]: #define the network
class Net(nn.Module):
    def __init__(self):
        super(Net, self).__init__()
        self.conv1 = nn.Conv2d(3, 6, 5)
        self.pool = nn.MaxPool2d(2, 2)
        self.conv2 = nn.Conv2d(6, 16, 5)
        self.fc1 = nn.Linear(16 * 5 * 5, 120)
        self.fc2 = nn.Linear(120, 84)
        self.fc3 = nn.Linear(84, 10)

    def forward(self, x):
        x = self.pool(F.relu(self.conv1(x)))
        x = self.pool(F.relu(self.conv2(x)))
        x = x.view(-1, 16 * 5 * 5)
        x = F.relu(self.fc1(x))
        x = F.relu(self.fc2(x))
        x = self.fc3(x)
        return x
```

```
In [5]: net = Net()
```

```
In [6]: #define loss function and optimizer
import torch.optim as optim
```

```
criterion = nn.CrossEntropyLoss()
optimizer = optim.SGD(net.parameters(), lr=0.001, momentum=0.9)
```

```
In [7]: #train the network with 2 epoch
        for epoch in range(2): # loop over the dataset multiple times

            running_loss = 0.0
            for i, data in enumerate(trainloader, 0):
                # get the inputs; data is a list of [inputs, labels]
                inputs, labels = data

                # zero the parameter gradients
                optimizer.zero_grad()

                # forward + backward + optimize
                outputs = net(inputs)
                loss = criterion(outputs, labels)
                loss.backward()
                optimizer.step()

                # print statistics
                running_loss += loss.item()
                if i % 2000 == 1999: # print every 2000 mini-batches
                    print('[%d, %5d] loss: %.3f' %
                          (epoch + 1, i + 1, running_loss / 2000))
                    running_loss = 0.0

            print('Finished Training')
```

```
[1, 2000] loss: 2.226
[1, 4000] loss: 1.916
[1, 6000] loss: 1.730
[1, 8000] loss: 1.619
[1, 10000] loss: 1.563
[1, 12000] loss: 1.478
[2, 2000] loss: 1.426
[2, 4000] loss: 1.371
[2, 6000] loss: 1.320
[2, 8000] loss: 1.316
[2, 10000] loss: 1.296
[2, 12000] loss: 1.266
Finished Training
```